


HeirFlow

Highly Available Workflow
Orchestration with Apache Airflow

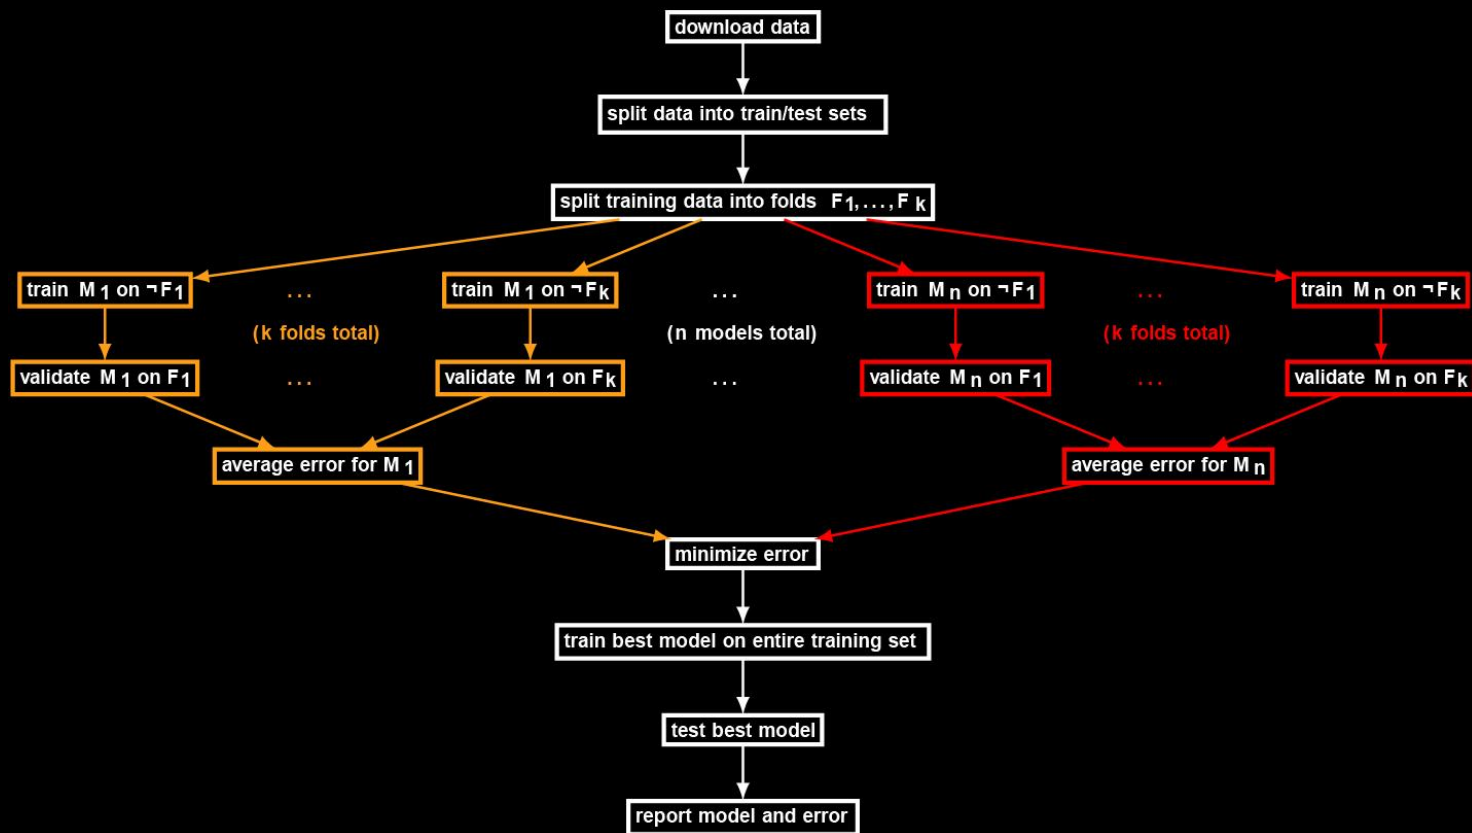
David Wiygul
Insight Data Engineering Fellowship, Los Angeles

Workflow

- Collection of interdependent tasks
- Represented as a directed acyclic graph (DAG)
- Vertices correspond to tasks
- Edges indicate dependence
- Often come with starting time and frequency
- Lifeblood of data engineering



Example DAG (machine learning)



Airflow components

- DAGs (Python code)




Airflow components

- DAGs (Python code)
- Metadata database




Airflow components

- DAGs (Python code)
- Metadata database
- Scheduler




Airflow components

- DAGs (Python code)
- Metadata database
- Scheduler
- Task queue



Airflow components

- DAGs (Python code)
- Metadata database
- Scheduler
- Task queue
- Worker

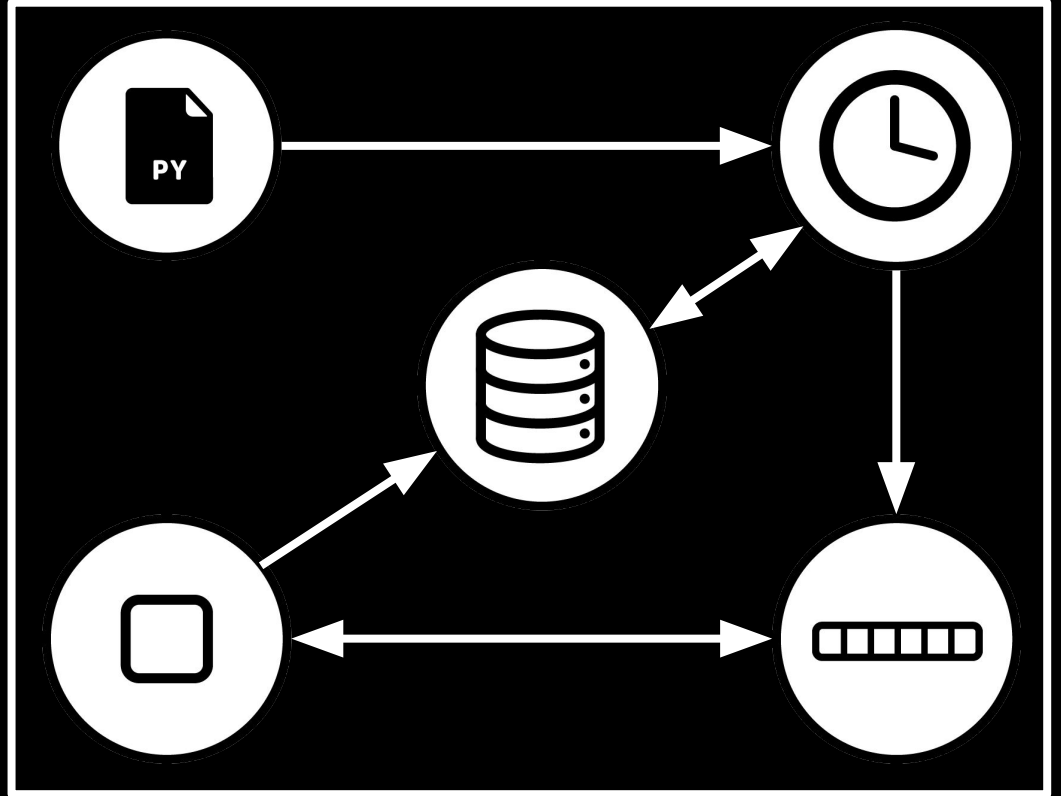




Airflow components

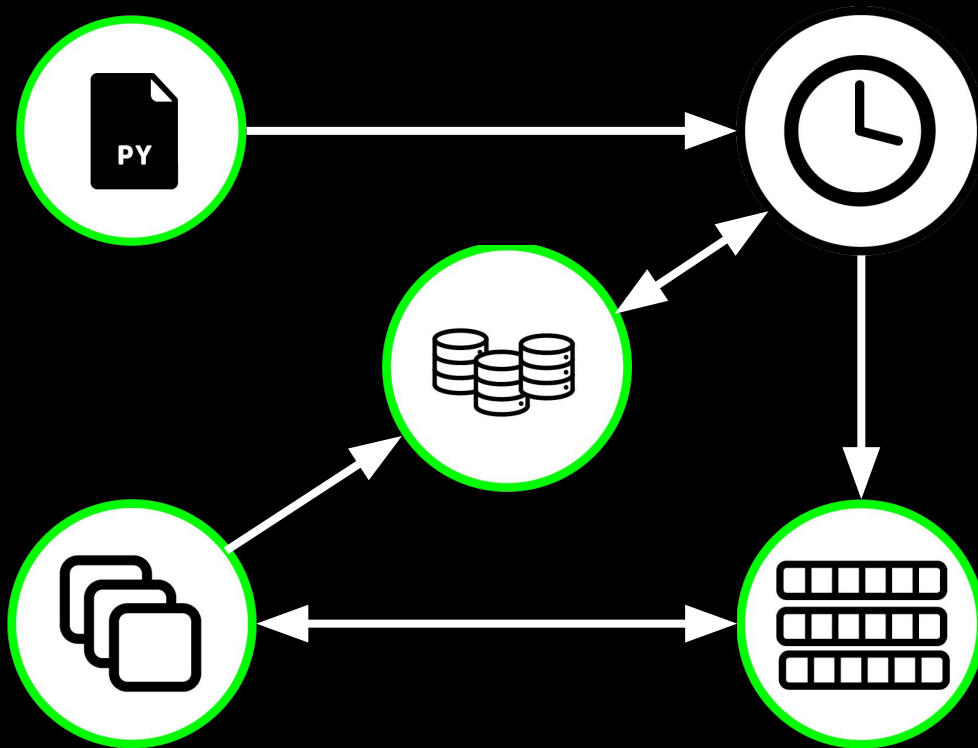
- DAGs (Python code)
- Metadata database
- Scheduler
- Task queue
- Worker

Default config: all on one machine



Horizontal scaling for availability and performance


- Downtime sucks.
- Interrupted workflows really suck.
- Parallelism rocks.
- Most components are straightforward to scale out.



Horizontal scaling for availability and performance

- Downtime sucks.
- Interrupted workflows really suck.
- Parallelism rocks.
- Most components are straightforward to scale out.

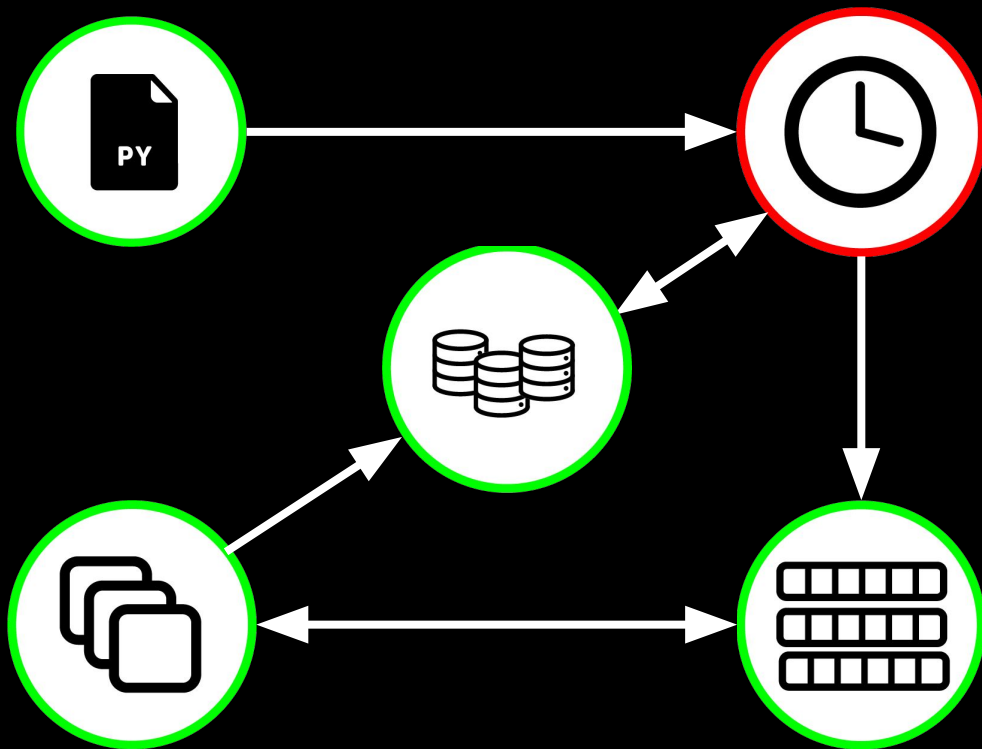
But not the
scheduler!



Scheduler redundancy

Schedulers running at the same time might schedule the same task repeatedly.

❌ Risky
❌ Wasteful



One solution

Use a cluster orchestration tool like Apache Zookeeper (or Consul) to ensure one scheduler is in action at one time

★ Convenient when already in place

✗ Otherwise requires additional resources and maintenance



Engineering challenge: scheduler redundancy

To ensure exactly one scheduler is running at a given time:

- Add a table (SSOT) to the already existing metadata DB.
- Wrap each scheduler in a Python script.
- The script checks in periodically with the DB.
- The oldest healthy scheduler takes charge.
- There is always a clear heir!

HeirFlow in action

File Edit

Scheduler
Scheduler
Scheduler
Scheduler
]


Airflow - DAGs


✕ +

← → ↻

Not secure | 52.38.71.0:8080/admin/




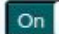









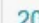

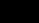
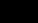
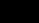
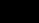
☆ 👤 ⋮

 Airflow



DAGs

Search:

		DAG	Schedule	Owner	Recent Tasks 	Last R
		model_select	1 day, 0:00:00	airflow	              	

GAME OF THRONES™



It's important to have a clear heir.

Pros and cons

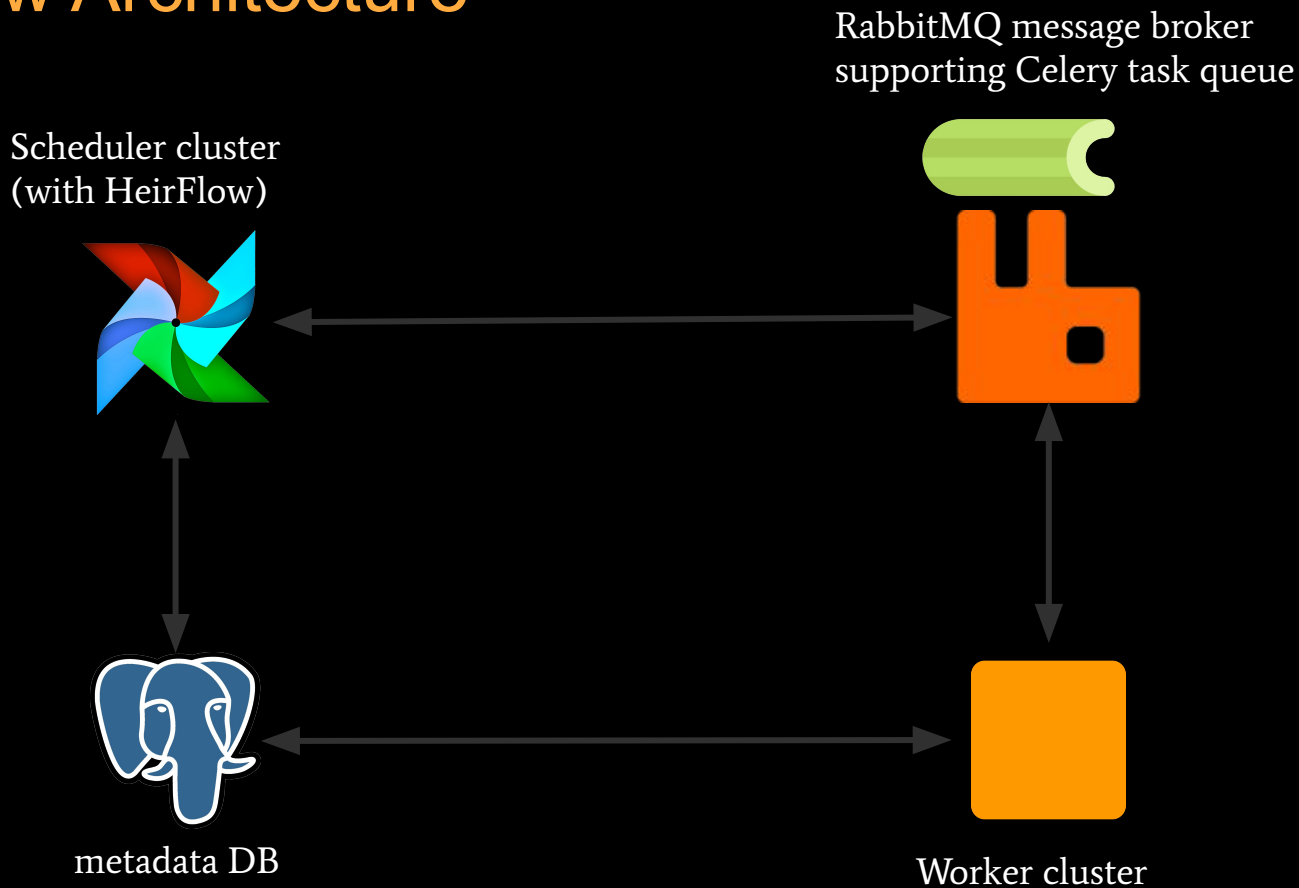
★ Availability

★ Lean setup

★ Flexibility

✗ Additional load on schedulers and (especially DB)

Airflow Architecture



About me: David Wiygul

- Before Insight I was a mathematician working in differential geometry and general relativity.
- Theme: complex systems obeying simple rules.
- As a data engineer I get to work with systems that offer (almost) instant feedback!
- I'm also proud dog-dad (see mutually recursive figure)...
- and an expecting human dad (yikes!)

